

Figure 1: Components of the DAQ

1 Purpose of this DAQ ?

I wrote this DAQ in the hope for it to become the (prototype of) standard DAQ used by IMP. As far as I learned, by the time this DAQ was written, no ‘standard’ DAQ existed in IMP and nobody was working on this, even (probably) nobody was thinking about such a thing. This is, in my humble opinion, one major thing that we were behind others (e.g. GSI, NSCL/FRIB). However, to really develop a universal DAQ applicable to many experiments, we may need a whole group to do that. I cannot do that on my own. So in this version of DAQ, I kept many things as simple as possible and many things remained unoptimized. The bottom line was to make sure that it works and can be easily used by others. How easy could it be ? Well, in most cases, one should be able to set it up by just clicking mouse and fill in some parameters (like module base addresses) without any coding (except for the online analysis part). To that end, one has to use only the modules predefined in this DAQ, unsupported modules won’t work properly (in fact, they won’t work at all). What if one needs to use an unsupported module? There are three ways to work it around: i) find an alternative module; ii) contact me to include your module (gaobsh@impcas.ac.cn); iii) do it yourself, you should be able to include your module easily since special care was taken to accomplish that.

Special care was also taken to make sure the DAQ can be upgraded to be used in more complicated experiments with minimal modifications (extensibility). That means although the current version is kept as simple as possible, it can be easily extended to a more complex version. Why not design it to be a complex DAQ in the first place? Well, there are many reasons. First, I don’t have time to do that. Second, more complicated system required high learning curve, which may hindrance people from using it. Third, as far as I can see, there is still no (or very little) such demands from the IMP.

Part of the DAQ designment is inspired by the NSCL DAQ.

2 Overview of the DAQ

This DAQ consists of the following parts, see Fig. 1:

- config.py. This is a GUI program (python script) to help user create the configuration file used by the DAQ.
- frontend. This is the part that communicates with the electronics (e.g. VME modules).

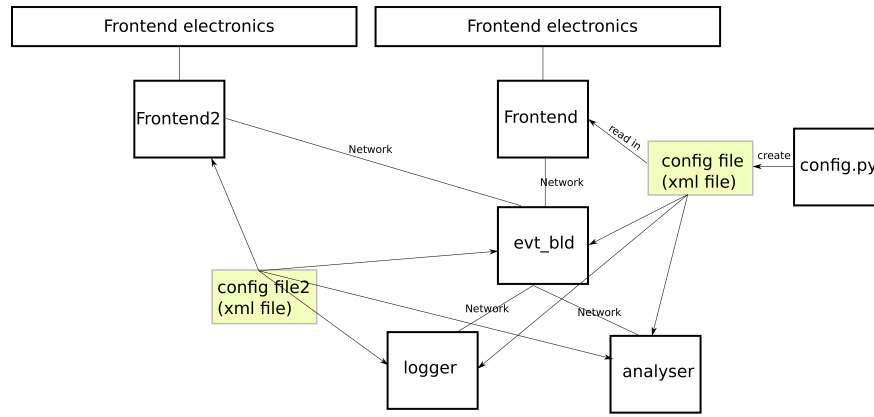


Figure 2: Configuration of the DAQ with two DAQs

- `evt_bld`. This is the event builder. It grabs data from the frontend and build a complete event based on timestamps of the fragments readout by frontend.
- `logger`. This program takes data from `evt_bld` and records it in hard drives.
- `analyser`. This program also takes data from `evt_bld`, it analysis the events and makes histograms instead of recording them.

The communications between different programs is done via sockets. I chose sockets instead of shared memory based on the following considerations: i) Different programs do NOT have to run in the same computer. This makes the DAQ more extendable. It also avoids the analyser slowing down the DAQ when it takes too much CPU time. ii) The synchronisation becomes easier because mutex or semaphores are not needed in between programs (although they may still be needed inside a program containing multi threads). Sockets are also easier to handle (as far as I'm concerned). iii) The sockets are of cause slower than shared memory, however, it should not be a problem in most cases.

In principle only the frontend and analyser may need the configuration, however, I designed it such that all components read in the configuration file at startup in case it may be needed in the future. When there are more than one DAQ needs to be used in an experiment, one can easily adjust the configuration as shown in Fig. 2. The only change is to change the `evt_bld` to be connected with two frontends.

3 Frontend

To minimize the dead time of DAQ, the buffers of the modules should be utilized. However, this imposes the problem of synchronisation. To simplify the DAQ, it is required that all supported modules should have a time stamp for each event. This might be a strong limitation. If one really wants to use their modules which don't have time stamps, they should design their electronics such that for each event, a busy signal is issued, which is reset after reading the event. This basically means the buffers of the modules are disabled. So people are strongly advised to use modules with time stamps. People may also need signals at the beginning and end of each run. So one requirement is that

after reading each event, a signal should be send out (for whatever purpose). At the beginning and end of each run, a signal should be send out. These three signals should use different ports.

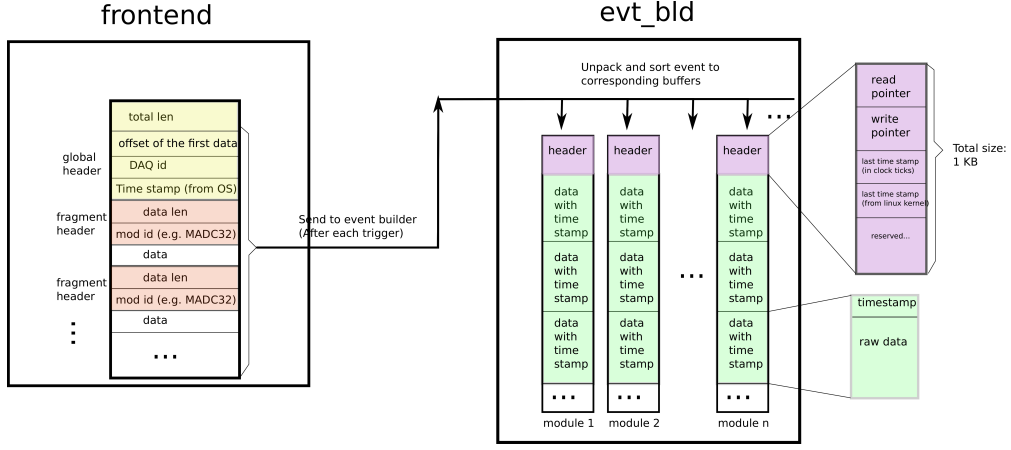


Figure 3: Data structures in frontend and event builder.

To maximize the speed,

The same type of modules should be readout in BLT mode (if they are in the same crate)

For example, if one has five MADC32 (in the same crate), they should be readout in BLT mode. Each BLT read may readout data for more than one event, if they are all send out directly to the event builder without preprocessing, the event builder will have to deal with a block of data from multiple events from multiple modules. It would be very difficult for the event builder to build events based on their time stamps (in principle it is still possible though). At the same time the preprocessing done by the frontend should be kept minimum to keep it fast enough. The final decision is like following:

In each BLT read, the data is packaged to contain the length of data (in byte) and the id of module where the data is readout. Then at the end of all readout, the whole data is packaged to contain the total length, the DAQ id, the offset of the first data and the time stamp (from the operating system). Then the whole package is send out to the event builder. See fig. 3.

The time stamp from the operating system is helpful to let the event builder determine the overflows of the time stamps of each module (see the section on event builder). The offset of the first data is the offset of the first fragment header within the whole data package. This offset should be used to determine the beginning of data. It allows variable length of the global header. The DAQ id may be useful if more than 1 DAQs are present. To maximize the speed, at least two thread should be used in the frontend. One reads out data from modules and packs it, the other sends the data out to event builder via network. The two threads use ring buffer to shared data and sequence locks for synchronisation. The structure of the ring buffer is shown in Fig. 3.

4 event builder

The main task of event builder is to build event from the data send from frontend based on their time stamps. To make the building process easier, the following methods are employed:

After receiving each data package from frontend, the data are sorted and saved to different buffers based on their origin.

Here each module has a buffer, not each type of module has a buffer. For example, one has 5 MADC32, then 5 buffers are created, one for each MADC. This is not too crazy. Imagine one has 100 modules for an experiment, there are 100 buffers. However, each buffer doesn't have to be very large. So the total memory consumption should not be a problem. And this also makes the building event process easier.

In each buffer, the events should have been sorted based on their time stamps.

You can imagine this is also going to make the building process easier. The time stamps at the beginning of each data blocks should be monotonic (at least for each run). That means when sorting the events, the event builder should be able to detect the overflows of time stamps and correct for that. The header of each buffer has the value of 'the last time stamp' T_{last} and 'the last computer time' T'_{last} to help to achieve this. The T_{last} is in unit of the clock ticks and the T'_{last} is the time stamp added by the frontend (from linux kernel). The T_{last} is monotonic and never overflows. The 'computer time' added by frontend should also use the monotonic clock provided by the kernel. From the T_{last} , it is straight forward to calculate the last time stamp readout from module $t_{last} = (T_{last} \bmod R)$ where R is the range. For example, if one module has a time stamp of 40-bit length, the range is $R = 2^{40}$. Since each event has the time stamp read out from module t_{this} and time stamp added by frontend T'_{this} , we can use these information together with those from the buffer header to get the following three quantities: $\Delta T'$, t_{last} and t_{this} which are the time interval between the two events, last time stamp readout from module and the current time stamp readout from module. These three quantities should have enough information to let us determine the number of overflows between the two events (if there are any).

How to build events?

Since each event in the buffers are already sorted by time, it is straight forward to combine them into one events. The tricky thing is when to start and stop to build events? The most simple answer would be to start building events immediately after receiving data from frontend and stop it when no more data in the buffers. However, this approach may have some problems when multiple events are readout in one data blocks. This is also shown in Fig. 4. The data with time stamp t_4 from module m and module n should belong to one event, however, they won't be combined together to form one event. A better way is to start building events whenever received data from frontend, as before. However, we stop building events as soon as any of the event buffers is empty. But there is one exception: if the buffer is empty at the beginning of event building, we should not stop. That means if we find out that one event buffer is empty before we retrieve any data from it, it just means the corresponding channel is not fired, we shouldn't stop our event building process. To recap:

We start event building whenever we received data from frontend, we stop event building if *any* of the buffers is empty. However, if one event buffer is already empty before we get any data out of it, we won't stop event building even if this buffer is empty.

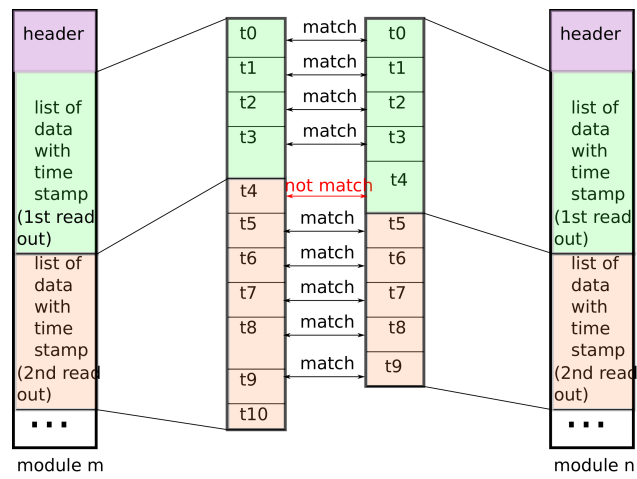


Figure 4: Data structures in frontend and event builder.